

SWIN
BUR
* NE *

SWINBURNE
UNIVERSITY OF
TECHNOLOGY

COS20031

Computing Technology Design Project

Week 09
Indexes





Database Development Lifecycle

1. Planning
2. Requirement gathering
3. Conceptual design
4. Logical design
5. Physical design

6. Construction

7. Implementation & rollout

8. Ongoing support



1. What & Why index?
2. Creating an index
3. Showing indexes
4. Dropping indexes
5. Multicolumn indexes
 - a. performance experiment
6. Project update



(A) What & Why indexes




What is index?

- Index is a data structure that helps (esp. in large data set)
 - rapidly look up data (WHERE clause)
 - enforce constraints
 - sort data (ORDER BY clause)
- Costs associated with index:
 - performance cost: needs updating when data is updated (insert, update, drop)
 - storage cost
 - should only be used with selected columns
- Common index design: B-Tree (Binary tree)
- Indexes are dropped when the associated table is



Binary Tree (B-Tree)

Find 145?





When to use index

- Primary key
- Columns used for ORDER BY
- Columns used in WHERE clauses:
 - col = value
 - col > value



(B) Creating an index

Creating an index



- Add to the CREATE TABLE statement
 - can be UNIQUE or Non-UNIQUE

```
CREATE TABLE test (  
    id INT UNSIGNED NOT NULL AUTO_INCREMENT PRIMARY KEY,  
    string1 VARCHAR(128),  
    string2 VARCHAR(128),  
    INDEX str1(string1)  
);
```

explicit index on
column (non-unique)

implicit index on
PK (also unique)



SHOW INDEX FROM test;

unique-
ness

index
name

data
structure

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
test	0	PRIMARY	1	id	A	0	NULL	NULL		BTREE
test	1	str1	1	string1	A	0	NULL	NULL	YES	BTREE



```
CREATE TABLE test (  
    id INT UNSIGNED NOT NULL AUTO_INCREMENT PRIMARY KEY,  
    string1 VARCHAR(128) UNIQUE,  
    string2 VARCHAR(128),  
    UNIQUE INDEX str1(string1)  
);  
  
SHOW INDEX FROM test;
```

UNIQUE indexes on
column

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
test	0	PRIMARY	1	id	A	0	NULL	NULL		BTREE
test	0	string1	1	string1	A	0	NULL	NULL	YES	BTREE
test	0	str1	1	string1	A	0	NULL	NULL	YES	BTREE

Creating an index (2)



- Use CREATE INDEX statement (after table creation)
 - can be UNIQUE or Non-UNIQUE

```
CREATE INDEX str2 ON test(string2);
```

OR

```
CREATE UNIQUE INDEX str2 ON test(string2);
```

```
SHOW INDEX FROM test;
```



(C) Showing indexes

Showing indexes (1): in one table



- Use SHOW INDEX FROM statement
- Variations:
 - SHOW INDEXES FROM...
 - SHOW INDEXES IN ...

Showing indexes (1): in all tables in a schema



```
SELECT DISTINCT table_name, index_name
FROM information_schema.statistics
WHERE table_schema = 'scratch';
```

table_name	index_name
customer	PRIMARY
item	PRIMARY
numerics	PRIMARY
sale	PRIMARY
test	PRIMARY
test	str2
test	str1



(D) Dropping indexes

Dropping indexes



```
DROP INDEX str1 ON test;
```

```
SHOW INDEX FROM test;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name	Collation	Cardinality	Sub_part	Packed	Null	Index_type
test	0	PRIMARY	1	id	A	0	NULL	NULL		BTREE
test	0	str2	1	string2	A	0	NULL	NULL	YES	BTREE



(D) Multicolumn indexes



- Index on more than one columns

```
CREATE TABLE test (  
    id INT UNSIGNED NOT NULL AUTO_INCREMENT PRIMARY KEY,  
    string1 VARCHAR(128),  
    string2 VARCHAR(128),  
    INDEX twostrs (string1,string2)  
);
```

can also add
UNIQUE

Multicolumn indexes



```
SHOW INDEX FROM test;
```

Table	Non_unique	Key_name	Seq_in_index	Column_name
test	0	PRIMARY	1	id
test	1	twostrs	1	string1
test	1	twostrs	2	string2

Multicolumn indexes: explanation



```
EXPLAIN SELECT string1, string2
FROM test
ORDER BY string1, string2;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	test	index	NULL	twostrs	774	NULL	1	Using index

Single column index: explanation



```
EXPLAIN SELECT string1  
FROM test  
ORDER BY string1;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	test	index	NULL	str1	387	NULL	1	Using index

Experiment: performance impact of index



- Drop table test
- Recreate it with the 2 column index
- Insert 45-50 records into the table
- Case 1: query with index (record exec. time)
 - use the 2-column SELECT query with ORDER BY on both columns (see a previous slide)
- Case 2: query without index
 - drop the index
 - execute the same query & record the exec. time
- Observe the execution time improvement.



(D) Project Update



- Incorporate indexes into your project
 - what columns, what indexes?
- Observe the effect of index
- Update Jira project



See Canvas.